

Step 1 – the correction is built from two small tables





Table A
trained, small

Table B
trained, small

Full-size correction
computed, never stored

Step 2 – added on top, not swapped in



Base weights
frozen, unchanged

LoRA correction
from Step 1

What the model reads
at inference